



Motivation & Method

Evaluate products, not isolated models.

Production work combines implicit constraints, heterogeneous documents and domain-specific acceptance criteria.

63%

of 27 surveyed AI-product companies report low confidence in model / prompt updates.

Two-layer construction framework

Human-led elicitation

Partner engagement + requirement elicitation

~1 month; 3–4 refinement cycles

Validated task package

query.md · task.yaml · files/

Fixed expert-authored .eval/rubric.py

Automated execution

Task Runner → Docker → evaluator dispatch

Record artifacts, trajectories and scores

Six occupational domains

O*NET domain	Tasks
Human Resources	11
Finance & Investment	22
Procurement & Operations	23
Software Engineering	11
Healthcare & Life Sciences	16
Technology Research	11

Human elicitation is not automated; the technical pipeline begins after expert validation.

Production Benchmark & Results





Figure 1. Research–production gap, requirement-to-benchmark construction, and main findings.



All 14 evaluated configurations

Model	Claude Code	Codex	GitHub Copilot	Cursor
Claude Opus 4.6	64.41	53.45	61.31	61.85
GPT-5.2	39.47	47.59	54.91	—
Gemini 3 Pro	50.78	—	49.92	—
Kimi K2.5	43.90	43.09	—	—
GLM-5	48.70	49.85	—	—
MiniMax M2.5	40.89	—	—	—

Overall score = equal-weight mean of six domain scores.
“—” denotes an untested pairing, not a zero score.

Large variation across domains

Mean score across the 14 configurations



Failure modes: cascade dependencies · subjective judgment · retrieval cross-section inconsistency · implicit constraints · output format